

2.4.1 Virtual Machines

- Infrastructure as a Service (IaaS)
 - Include a virtual processor, memory, storage, and networking resources
 - VM is a good choice when you need:
 - Total control over the OS
 - The ability to run custom software
 - To use custom hosting configurations
 - Azure takes care of the physical hardware
 - You take care of configuring, updating, and maintaining the software that runs on the VM
 - An **image** is a template used to create VM
 - Includes an OS and often other software, like development tools or web hosting environments
-

▼ Example Use Cases

- **During test + dev** as it's easy to create different OS and application configs. Easy to delete when not needed.
 - **Minor tasks** - Ex: application handles fluctuations in demand and shut down VMs when you don't need them & quickly start them to meet a suddenly increased demand
 - **Extending your datacenter** to the cloud - Ex: running Sharepoint
 - **During disaster recovery** - Ex: if primary datacenter fails, create VMs running on Azure to run your critical applications and then shut them down when the primary datacenter becomes operational again
 - **Lift and shift**: Moving from physical datacenter to cloud. You can take the image of the server & run within a VM with little to no changes
-

▼ Scaling and High Availability

- 99.99% uptime guarantee for all Virtual Machines that have two or more instances deployed across two or more Availability Zones

Domains & Maintenance Events

Update Domains

- Groups of VMs and underlying physical hardware that can be rebooted at the same time

Planned Maintenance Events

- When the underlying Azure fabric that hosts VMs is updated by Microsoft
- Done to patch security vulnerabilities, improve performance, and add or update features
- Often no impact, sometimes requires reboot
- When the VM is part of an availability set, the Azure fabric updates are sequenced so not all of the associated VMs are rebooted at the same time
 - VMs are put into different **update domains**

Fault Domains

- Fault domain = rack of servers
- Provides the physical separation of your workload across different power, cooling, and network hardware that supports the physical servers in the data center server racks
- In the event the hardware that supports a server rack becomes unavailable, only that rack of servers is affected by the outage

Unplanned Maintenance Events

- Involves a hardware failure in the data center like a power outage or disk failure
- VMs that are part of an availability set automatically switch to a working physical server so the VM continues to run

- The group of VMs that share common hardware are in the same fault domain

Availability Sets

- Logical grouping of two or more VMs that help keep your application available during planned or unplanned maintenance
- With an availability set, you get:
 - Up to three fault domains
 - each have a server rack with dedicated power and network resources
 - Five logical update domains
 - can be increased to a max of 20
- There's no cost for an availability set
 - Only pay for the VMs within the availability set
- Recommended for HA

VM Scale Sets

- Lets you create & manage a group of identical, load balanced VMs
- Allow you to centrally manage, configure, and update a large number of VMs to provide highly available applications
- The number of VM instances can automatically increase or decrease in response to demand or a defined schedule
- Helps you build large-scale services for areas such as compute, big data, and container workloads
- Provides high availability through regional or multiple Availability Zones deployment options

Azure Batch

- Large-scale job scheduling and compute management
- When running a job, batch:

- Starts a pool of compute VM for you
 - Installs applications and staging data
 - Runs jobs with as many tasks as you have
 - Identifies failures
 - Requeues work
 - Scales down the pool as work completes
 - Good for cases where you need raw computing power or supercomputer level compute power
-